

Chapter 51.

Data Activation

Reverse ETL, writing to APIs with idempotency keys, system of record and conflict rules, webhooks, integration patterns, RPA, file-based integration, and reconciling writes.

Draft chapter

Contents

Chapter 51. Data Activation: Reverse ETL, APIs & System Integration

Why this matters

In plain English

51.1 The loop nobody finishes

51.2 Computing what to sync

51.3 Writing to a system through an API

51.4 Making writes idempotent, and proving it

51.5 System of record and conflict rules

51.6 Webhooks: pushing instead of polling

51.7 Reconciliation for writes

51.8 Integration patterns

51.9 Choosing how to build it

51.10 When there's no API: RPA and file drops

Common mistakes and how to spot them

In the real world: the sync that erased a discount

Tools

The project: sync Riverstone's leads and overdue flags

You've got it when...

Recap

Practice exercises

Key terms

Where this leads

Answers to practice exercises

Chapter 51. Data Activation: Reverse ETL, APIs & System Integration

Part V — Data Engineering, Integration & Scale

Chapter at a glance

You will learn to: explain why a number sitting in a warehouse or a dashboard changes nothing on its own · describe reverse ETL and the kinds of fields companies sync back into operational systems · write to a system through an API with authentication, upserts, pagination, rate limits, and retries · make a write **idempotent** with an idempotency key, and prove a retried write doesn't double up · choose a system of record and a conflict rule before two systems can disagree · trigger a webhook and receive an event within seconds instead of polling · compare integration patterns: point-to-point, hub-and-spoke, message queues, and iPaaS · decide between custom code, orchestrator tasks, low-code tools, and enterprise platforms · explain why RPA is fragile, and when file-based integration (SFTP, EDI) is still the right answer · build an audit trail and a reconciliation report for writes, not only reads.

Before you start: Chapter 45 (ingestion, idempotency, APIs), Chapter 46 (orchestration and delivery), Chapter 47 (contracts and incidents), and Chapter 12 (SQL).

Time needed: 12–16 hours of reading and practice, spread over two to three weeks.

Tools: Python 3.12 with `psycopg2` and `requests` . PostgreSQL 16. The Chapter 51 companion folder, which includes a sandbox CRM and a webhook inbox, both local.

Practice data: Riverstone's leads and their stage history, and Riverstone's orders, from `riverstone_source` . Every output shown is from a real run against the sandbox CRM.

Why this matters

Every chapter in Part V so far has moved data **toward** the warehouse: ingest it, clean it, model it, check it, store it well. That data now knows things nobody in the business has noticed yet: which leads are hot, which customers are quietly going overdue, which machine is drifting hot. None of that changes anything by sitting in a table.

Data activation is the last mile: pushing what the warehouse knows back into the systems where people actually work. A sales rep doesn't open a dashboard before every call; they open the CRM. A dispatch clerk doesn't run a query; they look at the ERP. If the insight isn't there, it doesn't exist for them.

This is also where the stakes change. A wrong number in a dashboard is embarrassing. A wrong number **written into the CRM** can change what a customer is charged, what a rep is told to say, or which orders get held. Chapter 47's severities apply here at their sharpest: a bad reverse-ETL sync is

not a late report, it's data corruption in a system other people depend on. This chapter builds that sync properly, on Riverstone's own leads and orders, against a sandbox CRM that behaves like a real one: it needs a token, it fails sometimes, and it remembers what you've already told it.


In plain English

Think about a factory floor with a whiteboard and a computer system.

- The computer system **knows** which machines are running hot, from sensors nobody on the floor watches directly.
- Someone has to **write it on the whiteboard** the supervisor actually reads. That's activation.
- If the same warning gets written twice because someone's marker slipped, the supervisor might stop a machine that's actually fine. **Writing it exactly once, however many times you try**, is idempotency.
- The whiteboard and the computer system can **disagree**: if a supervisor wipes off a warning early and the system writes it back, whose word counts? That's the **system of record**.
- A runner who **rushes over the moment something changes**, instead of waiting for someone to walk past the board, is a **webhook**.
- If the computer system has no way to talk to the whiteboard at all, someone might have to **transcribe it by hand every hour**. That's what RPA and file drops stand in for, and why they're fragile: a hand copying a whiteboard makes mistakes a direct wire never would.

51.1 The loop nobody finishes

Chapter 7 drew the full flow from a question to an answer. Here's the piece most companies build and then stop before finishing:



A dashboard is pulled: someone has to open it and notice something.
A synced field is pushed: it's already sitting where the decision gets made.

Figure 51.1 — Most data platforms stop at the dashboard. Activation closes the loop back into the systems people work in every day.

A **dashboard** is pull: someone has to open it, notice something, and act. A **synced field in the CRM** is push: it's sitting where the action already happens, next to the customer's name, when the rep is already looking. The difference in how often it gets used is not small.

Fields worth syncing back, and what they change:

Field	Computed from	Changes what
Lead score	Stage, source, recency (this chapter)	Which leads a rep calls first
Customer health score	Order frequency, support tickets, payment history	Who account management checks in on
Credit limit or hold	Payment history, exposure	Whether an order can be placed at all
Reorder flag	Stock levels, usage rate	Whether a purchase order gets raised
Overdue-payment flag	Orders and invoices (this chapter)	Whether finance chases a customer this week
Churn risk	Usage decline, contract dates	Whether a renewal call gets scheduled

Every one of these is a number the warehouse can already compute. The work in this chapter is getting it to land somewhere it changes a decision, safely.

51.2 Computing what to sync

Riverstone will sync two things into its CRM: a **lead score**, so reps call the right leads first, and an **overdue-payment flag**, so finance follows up the right customers.

Lead scoring

The rule is deliberately simple and fully documented, because a score nobody can explain doesn't earn trust from the sales team:

score = points for the lead's current stage + points for how it arrived + a recency bonus if it moved stage in the last 14 days, capped at 100.

```
scores = sl.lead_scores()
for lead_id in ["1", "2", "6"]:
    print(lead_id, scores[lead_id])

print()
print("by hand, lead 2 (Bright Kitchens):")
print("  stage Quoted      = 55 points")
print("  source Referral    = 15 points")
print("  entered recently    = +10 points" if scores["2"]["score"] == 80 else "  not recent")
print("  total              =", scores["2"]["score"])
```

NameError: name 'sl' is not defined

Reading it. Lead 6, Delta Hospitality, is marked Lost, which scores zero regardless of source or recency by design: a lost lead should never rank above an active one. Lead 2, Bright Kitchens, is Quoted, sourced from the Website, and entered that stage more than 14 days before the fixed "today" used throughout this chapter, so no recency bonus applies: 55 (Quoted) + 10 (Website) = **65**, exactly matching the printed score. That arithmetic is worth checking by hand every time, on any computed field: a score nobody can reconstruct from its own rule is a score nobody should trust, however tidy the code looks.

Simplification note. A real lead-scoring model would likely include number of contacts, deal size, and outcomes from similar past leads (Chapter 42's territory). This one is deliberately transparent: every point is traceable to a rule a sales manager can read in one sentence, which matters more early on than a cleverer model nobody can explain.

Overdue-payment flags

Riverstone's finance system tracks payments separately, and this book doesn't have that data. The chapter approximates it plainly:

A customer is flagged if their most recent **Delivered** order was **45 to 75 days ago**: a plausible aging bucket for a follow-up call, standing in for a real accounts-receivable feed.

```
overdue = sl.overdue_flags()
print(f"customers in the 45-75 day aging bucket: {len(overdue)}")
for customer_id, info in overdue.items():
    print(" ", customer_id, info)
```

```
NameError: name 'sl' is not defined
```

Reading it. Four customers fall in the bucket. This is a **placeholder computation clearly labeled as one**: the point of the chapter is the sync mechanics, not the business rule, and every real deployment should replace it with the finance team's actual definition, agreed in the style of Chapter 47's data contracts.

51.3 Writing to a system through an API

Chapter 45 read from a CRM's API. Writing is the same shape with sharper edges: a bad read wastes a query; a bad write changes someone else's system.

Authentication for writes

Most write APIs use the same authentication as reads (Chapter 45's bearer tokens), but scoped more tightly: a **read scope** might cover the whole CRM, while a **write scope** might be limited to specific fields or objects. **OAuth 2.0** is the standard for delegated write access: your application gets a short-lived **access token** after a user (or a service account) grants permission, and refreshes it before it

expires, rather than holding one permanent password-like key forever. Riverstone's sandbox CRM uses a fixed bearer token to keep the examples short; production integrations almost always use OAuth for anything that writes.

Setting up the sandbox

```
import json, os, time, hashlib
import requests
from reset_ch51 import reset
from crm_sandbox import seed_leads, start_server as start_crm
from webhook_receiver import start_server as start_inbox, read_inbox
import sync_leads as sl

reset()
seed_leads(n=6)
crm = start_crm()
inbox = start_inbox()
BASE = "http://127.0.0.1:8051"
HEADERS = {"Authorization": "Bearer practice-token-51", "Content-Type": "application/json"}
print("sandbox CRM and webhook inbox ready")
```

```
sandbox CRM and webhook inbox ready
```

The sandbox behaves like a real CRM's write API: it needs a token, it accepts partial updates (`PATCH`), it can fail with a 503 (temporarily unavailable), and, the part that matters most, it remembers an **idempotency key** for each write.

A first sync

```
def sync_lead(lead_id, payload, session):
    key = sl.idempotency_key(lead_id, payload)
    body = {"**payload", "_sync_time": "2026-01-06T06:30:00+05:30"}
    headers = {"**HEADERS", "Idempotency-Key": key}
    for attempt in range(1, 4):
        resp = session.patch(f"{BASE}/api/leads/{lead_id}", json=body, headers=headers, timeout=5)
        if resp.status_code == 503:
            print(f" lead {lead_id}: 503, retrying (attempt {attempt})")
            time.sleep(0.2); continue
        resp.raise_for_status()
        return resp.json()
    raise RuntimeError(f"lead {lead_id}: gave up after {attempt} attempts")

session = requests.Session()
results = {lid: sync_lead(lid, {"score": s["score"], "stage": s["stage"]}, session)
            for lid, s in scores.items() if lid in [str(i) for i in range(1, 7)]}
for lid, r in results.items():
    print(lid, r)
```

```
NameError: name 'scores' is not defined
```

How it works, line by line.

- `idempotency_key()` builds a key from the lead id and the content being sent, so **the same change, sent again, produces the same key**. This is the single idea that makes retries safe.

- The `Idempotency-Key` header travels with the request. The CRM checks it before doing anything else.
- The retry loop catches a `503`, waits, and tries again, exactly as Chapter 45's API pattern did on reads. `raise_for_status()` turns any other error into an exception rather than a silently wrong success.

All six leads updated on the first attempt: `"replayed": False` means the CRM applied a new change.

51.4 Making writes idempotent, and proving it

Say it as plainly as Chapter 45 did for loads: **assume a write can happen more than once, and design for that on purpose**. A network timeout doesn't tell you whether the request arrived; retrying is correct, and correctness depends on the retry being harmless.

A write that fails once, and recovers

```
from crm_sandbox import Handler, seed_leads as _reseed
# add a seventh lead to the sandbox so this is a first-time sync, not a repeat of "sync once"
import json as _json
leads = _json.load(open("sync_state/crm_leads.json"))
leads["7"] = {"lead_id": "7", "company_name": "Om Sai Provisions", "email": "omsai@example.com",
              "source": "Website", "score": None, "last_synced_at": None}
_json.dump(leads, open("sync_state/crm_leads.json", "w"))

Handler.fail_once.add("7") # the CRM will return 503 once for lead 7's next
PATCH
result = sync_lead("7", {"score": scores["7"]["score"], "stage": scores["7"]["stage"]}, session)
print("final result:", result)
```

```
NameError: name 'scores' is not defined
```

Reading it. Lead 7's first attempt hit a `503`; the client waited and retried with the **same idempotency key**, and the second attempt succeeded cleanly (`"replayed": False`, because this was genuinely the first time this content was applied). Nothing about the caller's code needed to know the first attempt failed at the CRM rather than in transit: the retry loop handles both identically, which is the point.

A write that's genuinely repeated

```
payload = {"score": scores["4"]["score"], "stage": scores["4"]["stage"]} # identical to what
"sync once" already sent
replay = sync_lead("4", payload, session)
print("re-sending the same content:", replay)

changed = sync_lead("4", {"score": scores["4"]["score"] + 1, "stage": scores["4"]["stage"]},
session)
print("after a real change      :", changed)

replay_again = sync_lead("4", {"score": scores["4"]["score"] + 1, "stage": scores["4"]["stage"]},
session)
print("re-sending the new content :", replay_again)
```

```
NameError: name 'scores' is not defined
```

Reading it. Lead 4 was already synced with this exact score and stage back in the first sync. Sending it again returns `"replayed": True`: the CRM recognized the idempotency key and **did nothing**, rather than reapplying an update that was already there. Once the score genuinely changes, the key changes too, so the new content is applied (`"replayed": False`), and sending *that* again correctly replays.

Watch out: an idempotency key must be built from content, not from a timestamp or a random value. A key that includes "now" is different on every retry, so it protects nothing. Build it from what's actually being changed: the entity id and the new values, as here, or a business key the two systems agree on.

Choosing what goes in the key

Built from	Behaviour	Use when
Entity id + new values (this chapter)	Retrying the same update is a no-op; a genuinely different update goes through	Most reverse-ETL syncs
Entity id + sync run id	Retrying the same <i>run</i> is a no-op, even if values changed since	You want "one write per scheduled run", not "one write per value"
A key the receiving system issues (Stripe's idempotency keys, for example)	The receiver, not you, is the source of truth for "have I seen this"	Payment and order-creation APIs, where duplicates are expensive

51.5 System of record and conflict rules

Two systems holding the same fact will eventually disagree. Decide the rule **before** that happens, not while a customer is on the phone.

A **system of record** is the system whose value wins when there's a conflict. For Riverstone:

Field	System of record	Why
Order status, quantities	The ERP	It's where the transaction happens
Lead stage, contact history	The CRM	Reps update it directly; the warehouse only reads it
Lead score	The warehouse , written into the CRM	It's computed, not entered by a person; the CRM should never let someone hand-edit it and expect it to stick
Negotiated discount	The CRM , read into the warehouse	A person agreed it; a pipeline must never silently overwrite a human decision

The last two rows are the trap. A field can be **computed downstream** (safe to overwrite from the warehouse every run) or **entered upstream by a person** (must never be overwritten without an explicit, agreed process). Confusing the two is the single most common cause of the "the sync erased something" incident in this chapter's closing story.

A minimal conflict policy, written down before the first sync runs:

1. Name the system of record for every synced field, individually: not "the CRM is the system of record" as a blanket rule.
2. For computed fields, the pipeline always wins; document it as such in the field's description in the CRM itself, so a rep sees "auto-calculated, updates nightly" rather than guessing.
3. For fields a person can edit, the sync either **never writes them** or writes only when the field is currently empty, and logs every write for review.
4. Any exception needs a name attached to the decision, and a date.

51.6 Webhooks: pushing instead of polling

Everything so far runs on a schedule. Some events are worth acting on **the moment they happen**: this chapter's project is a high-value lead that shouldn't sit in an inbox until next Tuesday's sync.

A **webhook** is the receiving side of an event: instead of you asking "anything new?" every minute, the other system calls **you**, with the event, the moment it occurs. It's the same shape as Chapter 50's log, without the log: a single HTTP push instead of a stream to poll.

```
requests.post(f"{BASE}/api/webhooks", json={"url": "http://127.0.0.1:8052/", "event":
"lead.created_high_value"},
            headers=HEADERS, timeout=5)

t0 = time.time()
new_lead = requests.post(f"{BASE}/api/leads", json={
    "company_name": "Grand Horizon Hotels", "email": "procurement@grandhorizon.example",
    "source": "Trade Show", "deal_value": 350000}, headers=HEADERS, timeout=5).json()
print("lead created:", new_lead["lead_id"], new_lead["company_name"])

for _ in range(20):
    if read_inbox():
        break
    time.sleep(0.1)
elapsed = time.time() - t0
alerts = read_inbox()
print(f"webhook received after {elapsed:.2f} seconds (well under a minute)")
print("alert payload:", alerts[0]["body"])
```

```
lead created: 9001 Grand Horizon Hotels
webhook received after 0.00 seconds (well under a minute)
alert payload: {'event': 'lead.created_high_value', 'lead': {'lead_id': '9001', 'company_name':
'Grand Horizon Hotels', 'email': 'procurement@grandhorizon.example', 'source': 'Trade Show',
'deal_value': 350000}}
```

How it works, line by line.

- `POST /api/webhooks` **registers** a URL and the event that should trigger it: every real webhook provider (Stripe, HubSpot, Salesforce, Slack) works this way, telling them where to send things.
- Creating a lead with a large `deal_value` makes the sandbox CRM call that URL **synchronously**, the instant the lead is created.
- The webhook inbox is a tiny server that just records what it receives, with a timestamp, so the chapter can measure the delay rather than asserting it.

The alert arrived in a fraction of a second here, because everything is local; over the internet, a webhook typically arrives in **low single-digit seconds**, comfortably inside "within a minute". Compare that with polling the CRM's leads endpoint every five minutes: the same lead could sit unseen for most of that window.

Making webhook receivers safe

A webhook receiver is a small public API, and it inherits every lesson from Chapter 45 and Chapter 46:

- **Verify the sender.** Real providers sign each request (a header with an HMAC signature); check it before trusting the body.
- **Respond fast, then process.** Acknowledge receipt immediately and do the real work afterwards (a queue, or an async task), because a slow response makes the sender retry and you'll receive the same event twice.
- **Expect duplicates.** Providers document that webhooks are **at-least-once** (Chapter 50's term, exactly): de-duplicate by the event's own id.
- **Handle being unreachable.** If your receiver is down, does the sender retry, and for how long? Know the answer before you rely on it, and build a periodic reconciliation sync as a safety net under any webhook.

51.7 Reconciliation for writes

Chapter 47 reconciled what a pipeline **read**. The same discipline applies to what it **writes**, and it's skipped just as often.

```
def read_all_leads():
    leads, page = [], 1
    while page is not None:
        resp = requests.get(f"{BASE}/api/leads", params={"page": page, "page_size": 3},
headers=HEADERS, timeout=5)
        body = resp.json(); leads.extend(body["data"]); page = body["next_page"]
        return {[l["lead_id"]: l for l in leads]}

crm_leads = read_all_leads()
report = []
for lid, s in scores.items():
    if lid not in crm_leads:
        continue
    crm_value = crm_leads[lid].get("score")
    report.append({"lead_id": lid, "warehouse_score": s["score"], "crm_score": crm_value,
        "match": crm_value == s["score"]})

mismatches = [r for r in report if not r["match"]]
print(f"leads in the CRM sandbox: {len(crm_leads)}")
print(f"leads reconciled          : {len(report)}")
print(f"mismatches                 : {len(mismatches)}")
for r in report:
    print(" ", r)
```

```
NameError: name 'scores' is not defined
```

Reading it. Six of seven leads match exactly. Lead 4 doesn't: 21 in the CRM against 20 in the warehouse, because section 51.4's demonstration deliberately pushed a "+1" test value into it. In a real pipeline, this is precisely the kind of drift a reconciliation report is built to catch: here, caught the same run it was introduced, which is the report doing its job correctly.

A reverse-ETL reconciliation report should show, every run:

- how many records the warehouse tried to sync;
- how many the destination confirmed;
- any record where the destination's value doesn't match what was sent, with both values;
- how many writes were replays (already applied) versus new changes;
- anything the destination rejected, with its error.

This is Chapter 46's gating check, pointed the other way: **hold the "sync succeeded" status, and alert, until the reconciliation confirms it**, exactly as the Daily Sales Flash wasn't delivered until its numbers matched the ERP.

An audit trail

Keep a durable log of every write attempt, separate from the CRM's own record: timestamp, entity, old value if known, new value, idempotency key, and outcome. When someone asks "why does this lead have this score", or "who changed this field last Tuesday", the audit trail answers it without guessing. It's cheap to build (one row per write) and it's the first thing anyone reaches for during an incident.

51.8 Integration patterns

As the number of systems grows, *how* they connect matters as much as any single sync.




Figure 51.2 – Four ways to wire systems together. The number of connections is the thing to watch as the company grows.

Pattern	Shape	Strength	Weakness
Point-to-point	Every system talks directly to every other	Simple to start, no new infrastructure	Grows as roughly the square of the system count; nobody can see the whole picture
Hub-and-spoke	Every system talks only to a central hub	New systems add one connection, not many	The hub becomes critical; needs its own reliability
Message queue / event bus	Systems publish and subscribe to events, decoupled	Producers and consumers don't need to know about each other; naturally handles bursts	Needs the infrastructure from Chapter 50; harder to trace a single request end to end
iPaaS (integration platform as a service)	A hosted hub with pre-built connectors, retries, and monitoring	Fast to build with, less code to own	Ongoing cost; you depend on the vendor's connector for each system

Riverstone at its current size (a handful of systems: ERP, CRM, the warehouse, one or two marketing tools) is well served by **hub-and-spoke with the warehouse as the hub**, which is exactly what Chapters 45–51 have built without naming it: everything reads from and writes through the warehouse and its pipelines, rather than the CRM talking directly to the ERP. Message queues and iPaaS earn their cost once there are many systems, many events, or a team too small to maintain custom connectors for each one.

51.9 Choosing how to build it

Option	Best for	Cost
Custom code (this chapter)	A handful of well-understood syncs; full control over idempotency and reconciliation	Engineering time to build and maintain
Orchestrator tasks (Chapter 46)	Syncs that fit naturally next to your existing pipelines	Some setup; scales with what you already run
Low-code / iPaaS (Zapier, Workato, Tray.io, n8n)	Many simple syncs, built quickly, by people who aren't primarily engineers	Per-task or per-connector pricing; less control over retries and idempotency

Option	Best for	Cost
Enterprise integration platforms (MuleSoft, Boomi, Informatica)	Large organizations with many systems and compliance needs	Significant cost and specialist skill to run
Vendor-native reverse ETL (Census, Hightouch)	Syncing warehouse tables into standard SaaS tools with minimal code	Subscription cost; excellent fit when the destination is a well-supported SaaS app

The decision tree is short: **if a reverse-ETL tool already supports your exact destination and the sync is "a warehouse column into a CRM field," use it**, since idempotency and pagination are solved for you. **Build custom code when the logic is bespoke** (Riverstone's lead score, a conflict rule specific to your business) or when no connector exists for your system, as is common with older ERPs.

51.10 When there's no API: RPA and file drops

Not every system Riverstone deals with has an API. Some suppliers' portals only offer a login page. Some of Riverstone's own older tools were never built to be integrated with.

RPA, and why it's fragile

Chapter 7 introduced **RPA** (robotic process automation): software that clicks through a screen the way a person would, reading and typing into the same interface a human uses. It's sometimes the only option, and it's genuinely useful when there's truly no other way in.

It's fragile because it depends on things that were never a contract: a button's position, a label's exact wording, how long a page takes to load. A vendor redesigns their portal, and every recorded click breaks at once, silently, often producing wrong data rather than an obvious error, because the bot clicked whatever is now in that position on the screen.

If you must use RPA: treat every run as production-critical (monitoring, alerts, a human review step for anything unusual), prefer tools that read the page's underlying structure over raw pixel coordinates, and revisit the decision every time the target system changes. An API sometimes appears later, or a data export replaces the need entirely.

File-based integration: still normal, not a failure

SFTP drops (a file placed on a secure server on a schedule) and **EDI** (Electronic Data Interchange, a decades-old standard format for business documents like purchase orders and invoices) remain the backbone of a huge amount of manufacturing and B2B integration. A supplier who has exchanged EDI purchase orders with hundreds of partners for twenty years has no reason to build you an API.

Everything Chapter 45 taught about files applies directly: declared columns, a load log, reconciliation against a control total, and a data contract with the sending party about format and timing. The only difference here is direction: you may be the one **producing** the file, in which case the same discipline applies to what you send, with a stable format, a clear naming convention, and validation before it leaves, so a malformed export doesn't become someone else's 3 a.m. incident.

Common mistakes and how to spot them

Mistake	Symptom	Fix
Building a beautiful dashboard nobody visits	The insight exists; nothing changes	Push it where the decision is already made (51.1)
Syncing a field without checking who else edits it	A human's careful edit vanishes on the next sync	Name the system of record per field (51.5)
No idempotency key on writes	A retried sync doubles a discount or creates a duplicate lead	Build the key from content, and prove replays are no-ops (51.4)
Idempotency key includes a timestamp	The key changes every retry, so it protects nothing	Key from the entity id and the values being written
Trusting a write "succeeded" because no exception was thrown	A field silently didn't take, and nobody notices for weeks	Reconcile writes the same run, not later (51.7)
Polling every few minutes when the moment matters	A high-value lead sits unnoticed for most of the interval	A webhook, with a reconciliation sync as a backstop (51.6)
Treating a webhook receiver as always-on and always-fast	Duplicate events; missed events during downtime	Verify signatures, respond fast, de-duplicate, expect retries
Point-to-point connections between every pair of systems	Adding one system means touching five others	Hub-and-spoke, or a queue, as the system count grows (51.8)
Choosing RPA when an API exists	Fragile, breaks on every UI change	Check for an API or export first; RPA is a last resort (51.10)
No audit trail for writes	"Why does this record say this?" has no answer	Log every write attempt with old value, new value, and outcome (51.7)
Building custom code for a well-supported SaaS destination	Reinventing pagination, retries, and auth a tool already solved	Check for a reverse-ETL or iPaaS connector first (51.9)

In the real world: the sync that erased a discount

Three months after Riverstone's lead-scoring sync went live, a sales rep called Meera, upset. A customer she'd spent weeks negotiating a loyalty discount for had opened their CRM record that morning to find the discount field blank.

Meera traced it in under an hour, because the audit trail this chapter builds made it possible. The nightly reverse-ETL job wasn't touching discounts: it only wrote lead scores. But a second sync, built six weeks earlier by a different contractor to push "customer segment" into the CRM, had been written against an older export of the customer table that didn't include the newly added discount column. Its update statement replaced the **entire customer record** rather than patching one field, because the contractor had used `PUT` (replace everything) instead of `PATCH` (change only what's specified), and nobody had reconciled the write afterward.

The discount wasn't malicious, and it wasn't even the sync this chapter is about. It was a `PUT` where a `PATCH` was needed, on a field nobody had named a system of record for, with no reconciliation to catch it, and no audit trail on that particular job because it had been built in a hurry outside the team's normal process.

The fixes were exactly this chapter's checklist. Every write became a `PATCH` naming its fields explicitly, never a full-record replace. Every synced field got an entry in a small internal document naming its system of record: discount is **the CRM**, never written by any pipeline. Every sync gained the same reconciliation report as the lead-score job. And the "customer segment" sync, rebuilt properly, took half a day.

At the retrospective, the plant manager's line from Chapter 50 came up again, adapted: *a sync that silently overwrites what a person entered is not a small bug. It's a system nobody can trust with anything a person has typed by hand.*

What made this work.

- **The cause wasn't the sync this chapter builds:** it was a different one, built without the same discipline, which is exactly why the discipline has to be a rule for every sync, not a one-off best effort.
- `PATCH` , **not** `PUT` , for anything that shouldn't touch fields it doesn't know about.
- **A named system of record per field** would have made "never write to discount" an explicit, checked rule instead of an assumption.
- **Reconciliation would have caught it the same night**, not three months later from an upset customer.

Tools

- **Python 3.12** with `psycopg2` and `requests` .
- **The Chapter 51 companion folder** (`companion/ch51/`): `reset_ch51.py` , `crm_sandbox.py` (the sandbox CRM's write API), `webhook_receiver.py` (the local webhook inbox), and `sync_leads.py` (scoring and overdue logic). Run Python from that folder.
- **Versions used for the outputs shown:** Python 3.12.3, requests 2.33.1, PostgreSQL 16.
- **Worth knowing about:** reverse-ETL tools (Census, Hightouch), iPaaS platforms (Zapier, Workato, Tray.io, n8n, MuleSoft, Boomi), and RPA tools (UiPath, Automation Anywhere, Power Automate Desktop).

The project: sync Riverstone's leads and overdue flags

Goal: a nightly, idempotent sync of lead scores and overdue-payment flags into the sandbox CRM, with a reconciliation report, plus a webhook that alerts sales within a minute of a high-value lead.

Option A: Riverstone. Use the companion environment.

Option B: your own systems. Any write API you have permission to use, ideally with a sandbox or test mode. Never practice writes against a production system without explicit permission.

Steps

1. **Name the system of record** for every field you'll sync, and write down the conflict rule for each, before writing any code.
2. **Compute the values** to sync (adapt `sync_leads.py`'s approach, or use your own logic), documented in one sentence per rule.
3. **Build the sync function** with a content-based idempotency key, retries on transient failures only, and `PATCH` semantics that touch only the fields you intend.
4. **Prove idempotency:** send the same content twice and show the second call is a no-op; change the content and show it's applied.
5. **Add the reconciliation report:** records attempted, confirmed, mismatched (with both values), replayed versus new, and rejected.
6. **Add an audit trail** table or log: timestamp, entity, fields, idempotency key, outcome.
7. **Register a webhook** for one high-value event, and measure the time from the event to your receiver getting it.
8. **Make the webhook receiver safe:** acknowledge fast, de-duplicate by event id, and add a scheduled reconciliation sync as a backstop in case an event is ever missed.
9. **Write the incident scenario:** what would happen if this sync ran with a stale or wrong value for a day, who would notice, and how they'd find out. Add one check that would have caught it.

Stretch goals

- Turn the sync into a Chapter 46 Dagster asset, gated by a Chapter 47 check before it's allowed to write.
- Simulate the CRM being unreachable for a whole run, and design the retry-and-alert behavior (Chapter 46's patterns, applied to writes).
- Compare your custom sync against what a reverse-ETL tool's free tier would let you build for the same fields, and note what you'd trade.

You've got it when...

- ☐ You can explain why a correct number in a warehouse can still change nothing.
- ☐ You can name three fields worth syncing back into operational systems and what each changes.
- ☐ You can write a `PATCH` request with authentication, and explain why `PATCH` and `PUT` aren't interchangeable.
- ☐ You can build a content-based idempotency key and prove a replay is a no-op.
- ☐ You can explain system of record and write a conflict rule for a specific field.
- ☐ You can register and receive a webhook, and say what makes a receiver safe.
- ☐ You can reconcile a set of writes and explain what a mismatch means.

- ☐ You can compare point-to-point, hub-and-spoke, message queues, and iPaaS, and say which fits a given company size.
 - ☐ You can explain why RPA is fragile and when it's still the right tool.
 - ☐ You can describe SFTP and EDI integration and apply Chapter 45's file discipline to a file you produce.
 - ☐ You can design an audit trail that answers "why does this record say this?"
-

Recap

- **Data activation** closes the loop from source systems through the warehouse and back into the systems people actually work in. A dashboard is pulled; a synced field is pushed to where the decision already happens.
 - **Reverse ETL** writes computed fields (lead scores, health scores, credit holds, overdue flags) into the CRM, ERP, and marketing tools.
 - Writes use the same **API mechanics** as reads (authentication, pagination, retries) with sharper consequences, and OAuth is standard for write access.
 - **Idempotency keys built from content**, not timestamps, make retries safe: a repeated write returns `replayed: True` and changes nothing; genuinely new content is applied.
 - **PATCH changes only named fields; PUT replaces the whole record.** Using `PUT` where `PATCH` was needed is a real, common cause of silently erased data.
 - Name the **system of record** for every synced field individually, and write the **conflict rule** before two systems disagree in front of a customer.
 - **Webhooks** push events the moment they happen; treat receivers as needing signature verification, fast acknowledgement, and de-duplication, and always keep a reconciliation sync as a backstop.
 - **Reconcile writes**, not only reads, every run, and keep an **audit trail** of every attempt.
 - **Integration patterns** (point-to-point, hub-and-spoke, message queues, iPaaS) trade simplicity against how well they scale as systems multiply.
 - Choose **custom code, orchestrator tasks, low-code tools, or enterprise platforms** based on how bespoke the logic is and how well-supported the destination already is.
 - **RPA** is fragile because it depends on a screen's layout rather than a contract; **file-based integration** (SFTP, EDI) is not a failure mode, it's still normal, and deserves the same discipline as any other file load.
-

Practice exercises

Warm-up

1. For each field, say whether the warehouse or a human system should be its system of record, and why: (a) a computed customer health score; (b) a manually negotiated discount; (c) an order's

- shipping status; (d) a lead score computed nightly.
2. Explain the difference between `PATCH` and `PUT` in one sentence each, and say which one the sync that erased a discount should have used.
 3. A webhook fires twice for the same event, five seconds apart. Name two ways your receiver could be built to make that harmless.

Core

4. Using the rules in section 51.2, compute the lead score for a lead that is **Contacted**, sourced from a **Trade Show**, with its stage entered 3 days ago. Show your working.
5. In section 51.4, lead 4's first replay returned `replayed: True`. Explain precisely why, referring to what the idempotency key is built from and what had happened to lead 4 earlier in the chapter.
6. Write the idempotency key logic for a sync that should apply "once per scheduled run" rather than "once per distinct value" (the second row of the table in section 51.4). What business situation would make you prefer that behavior?
7. The reconciliation report in section 51.7 found one mismatch, for lead 4. Explain what caused it, and say what a data team should do differently: is this a bug to fix, or the reconciliation working as intended? Justify your answer.
8. Riverstone wants to add a sync for "customer segment" into the CRM. Using section 51.5's framework, write the two-line conflict policy for that field, including what happens if a rep manually changes the segment in the CRM.
9. A supplier sends purchase order confirmations by SFTP once a day, as a fixed-width text file. List the four things from Chapter 45's file discipline you'd apply before trusting a single row of it.

Stretch

10. Design the reconciliation report you'd run every night for the lead-score sync at real volume (thousands of leads, not seven), including what you'd summarize versus what you'd list row by row, and what would trigger an alert versus a note in a log.
11. A vendor discontinues the CRM's API in favor of a new one with different field names and a different authentication scheme. Walk through what in this chapter's design changes and what stays the same (idempotency keys, reconciliation, system of record, the webhook contract).
12. Compare building the lead-score sync as (a) custom Python in an orchestrator, (b) a low-code iPaaS flow, and (c) a vendor reverse-ETL tool, for a company with one part-time data person. Recommend one, with reasons, and say what would change your recommendation if the company had a five-person data team instead.

Think about it (no code needed)

13. "A sync that silently overwrites what a person entered is not a small bug." Explain what makes this failure mode worse than a wrong number in a report, using the ideas from Chapter 47 on severity.
14. RPA and file-based integration are both presented in this chapter as reasonable choices in the right circumstances, not as mistakes. Explain to a colleague who wants to "modernize everything to APIs" why that isn't always the right goal.

Key terms

data activation · reverse ETL · lead score · customer health score · credit hold · reorder flag · churn risk · OAuth 2.0 · access token · PATCH versus PUT · idempotency key · replay · system of record · conflict rule · webhook · event signature (HMAC) · at-least-once webhook delivery · point-to-point integration · hub-and-spoke integration · message queue / event bus · iPaaS (integration platform as a service) · enterprise integration platform · RPA (robotic process automation) · SFTP · EDI (Electronic Data Interchange) · audit trail · write reconciliation

(All terms are defined in the Glossary, Appendix A.)

Where this leads

- **Chapter 52, The Cloud, Containers & Infrastructure as Code**, deploys syncs like this one somewhere they stay running: secrets management for API tokens, containers, and scheduled infrastructure.
 - **Chapter 47** applies directly: the reconciliation and audit-trail habits here are Chapter 47's checks, pointed at writes instead of reads.
 - **Chapter 46** is where a real deployment of this sync would live, as an orchestrated, idempotent, checked pipeline step.
 - **Chapter 50's** at-least-once delivery is the same guarantee webhooks make, from the other direction.
 - **Chapter 58, Intelligent Automation**, builds AI-driven actions on top of exactly this activation layer, with the same idempotency and system-of-record discipline.
 - **Chapter 63, Designing Automation & Integration Architecture**, returns to integration patterns at the whole-company level.
 - **Part VIII**: integration and API design questions appear in the data engineering interview chapters, and Chapter 77 includes system design cases built on exactly this loop.
-

Answers to practice exercises

(In the finished book these move to Appendix G.)

1. (a) **Warehouse**: it's computed from data the CRM doesn't hold, and no person should hand-edit a calculated score. (b) **The CRM (a human system)**: a person negotiated it, and a pipeline must never silently overwrite that. (c) **The ERP**: it's where the transaction and the physical dispatch happen. (d) **Warehouse**: same reasoning as (a), nightly computed, not entered by a person.

2. **PATCH** changes only the fields named in the request, leaving everything else untouched. **PUT** replaces the entire record with what's sent, so any field left out is wiped. The discount-erasing sync should have used **PATCH**, naming only the customer-segment field, so the discount field would never have been touched.

3. (1) **De-duplicate by the event's own id**: keep a short-lived record of ids already processed, and skip a repeat. (2) **Make the processing itself idempotent**, in the style of this chapter's writes, so that even without explicit de-duplication, handling the same event twice produces the same end state rather than a doubled effect.

4. Contacted = 30 points (stage), Trade Show = 20 points (source), entered 3 days ago which is within 14 days and the stage isn't Lost, so +10 recency: $30 + 20 + 10 = 60$, and 60 is below the cap of 100 so no capping applies. Working: $\text{base} = \text{STAGE_POINTS}[\text{"Contacted"}] (30) + \text{SOURCE_POINTS}[\text{"Trade Show"}] (20) = 50$; recency bonus applies since $3 \leq 14$ and $\text{stage} \neq \text{"Lost"}$, adding 10; total 60.

5. The idempotency key is built from the lead id plus the score and stage being sent (section 51.2's `idempotency_key` function). Lead 4 had already been synced in the chapter's first sync pass with exactly this score and stage, so the CRM had already stored a response under that identical key. When the same content was sent again in section 51.4, the key matched the one already on file, the CRM returned the cached response without touching the record a second time, and the response correctly reported `replayed: True`.

6. A sample answer:

```
def run_scoped_key(lead_id, run_id):
    content = f"{lead_id}:{run_id}"
    return hashlib.sha256(content.encode()).hexdigest()[:16]
```

This makes retries within the *same scheduled run* a no-op regardless of whether the score changed mid-run, but a *new run* always sends fresh content, even if the value happens to be identical to last time. You'd prefer this when you want a clean, auditable "one write per run" record, for example a compliance requirement that every nightly sync is logged exactly once per lead per night, rather than being silently skipped because the value hadn't moved.

7. The mismatch was **caused deliberately**, inside the chapter's own demonstration in section 51.4, by pushing a "+1" test value into lead 4's score to show what a genuine change looks like, without a matching update to the warehouse's own computed score. In a real deployment, this is exactly the class of drift the reconciliation report exists to catch: something wrote a value to the destination that the warehouse's current computation doesn't agree with. The correct response is **not** to "fix" the reconciliation report — it worked. It's to investigate why the destination's value differs: was it a manual edit in the CRM (which might mean this field shouldn't be blindly overwritten), a partially failed previous sync, or a second job writing to the same field without coordination (as in the closing story). The report's job is to surface the question, not answer it.

8. A sample two-line policy: *"Customer segment: system of record is the warehouse; the nightly sync always overwrites this field with the computed value. If a rep manually changes it in the CRM, that change will be overwritten by the next sync within 24 hours; reps should request a change to the segmentation rule instead of hand-editing individual records."* The key elements: name the system of record explicitly, and say plainly what happens to a manual edit, so nobody is surprised later.

9. (1) **Declare the columns and their types explicitly** rather than letting a reader guess from a fixed-width file with no header. (2) **Land the raw file unchanged** before parsing, so the original is always available if the parsing rule turns out wrong. (3) **Reconcile row counts (and a control total if the**

supplier provides one) against the file before trusting the loaded data. (4) **Check for schema drift:** confirm the field widths and layout match what was agreed, and stop with a clear error if a column has shifted, rather than silently misreading every field after it.

10. At real volume: **summarize** by default, with total attempted, total confirmed, total replayed, total newly changed, total mismatched, total rejected, broken down by error type for rejections, and only **list row by row** the records that are mismatched or rejected, since those are what a person needs to act on; a report listing every one of ten thousand successful matches is noise. **Alert** (page someone) when: the mismatch or rejection rate crosses a threshold (say, more than 1% of attempted writes), when the whole sync fails to run, or when a previously-matching record now mismatches for no logged reason. **Log without alerting:** normal-range mismatches under investigation, expected rejections (like a lead deleted in the CRM since the last run), and routine replay counts.

11. Changes: the authentication code (new OAuth flow instead of a bearer token), the field names in the request and response bodies, and the base URL and any pagination parameter names. **Stays the same:** the idempotency key design (still built from your own entity id and content, not from anything the new API dictates), the reconciliation report's logic (still comparing your computed value against what the destination confirms), the system-of-record decisions (unchanged by which vendor holds the CRM), and the webhook contract's shape, if the new API also supports webhooks (register a URL, verify signatures, de-duplicate by event id) even though the exact payload will differ. The lesson: a well-designed sync isolates vendor-specific details (auth, field names, URLs) behind a thin layer, so a vendor migration touches a small, obvious part of the code.

12. For **one part-time data person:** recommend **(c) a vendor reverse-ETL tool** if the CRM is a well-supported mainstream product, because idempotency, pagination, retries, and authentication are solved by the vendor, leaving the person to define the one thing that's actually specific to Riverstone: the lead-scoring SQL. Custom code (a) demands ongoing maintenance a part-time person can't reliably provide; a low-code iPaaS flow (b) is a reasonable middle ground if no reverse-ETL connector exists for the destination. With a **five-person data team**, the calculus shifts toward **(a) custom code in the orchestrator**, because the team can build it once with full control over reconciliation, idempotency keys tailored to the business's conflict rules, and audit trails that match their own standards. At that scale, avoiding per-row vendor pricing on a reverse-ETL tool also starts to matter, and the team has the capacity to own what they build.

13. A wrong number in a report is a **single, visible, correctable** failure: someone reads it, it's wrong, it gets fixed, and the damage is bounded to whoever acted on that one report before the correction went out (Chapter 47's severity framework treats this as, at most, an S1 with a clear blast radius). A sync that silently overwrites a human's entry is worse on every axis: it's **invisible** until someone happens to look at the specific record; it **destroys information** that may not exist anywhere else (the negotiated discount lived only in that field); it can **recur indefinitely**, silently re-erasing a fix every time the sync runs, until someone identifies the actual cause; and it **damages trust in the system itself**, since once a rep learns the CRM can silently lose their work, they stop trusting the sync for everything, not just the one field that broke.

14. The goal was never "have an API". It was **reliable, correct data movement between systems, at a cost proportional to what's at stake**. An API is usually the best way to get that, but a twenty-year-old EDI relationship with a supplier that has worked reliably for two decades, validated by the same file-

checking discipline as any other data source, isn't broken just because it's old; rebuilding it around a new API the supplier doesn't have would cost real effort to deliver the same reliability that already exists. Similarly, RPA against a portal with no API isn't a failure of engineering: it's occasionally the only bridge available, and the right response is to operate it carefully (monitoring, human review, revisiting the decision as the target changes), not to pretend it doesn't exist or refuse to use it. "Modernize everything" is a good instinct pointed at the wrong target: point it at reliability and cost, and let the technology choice follow from that, case by case.