

The product ladder, and the project catalogue

3 October 2026. Captured from Abhishek's design. Companion to [Selling-the-books-strategy-and-red-team](Selling-the-books-strategy-and-red-team.md), which covers pricing and go-to-market; this one covers what the products actually are.

The ladder

	Product	Price	What it is
1	The Interview Book	₹500	Book 4 alone. 827 questions today, heading for ~1,160. The thing Instagram points at.
2	The Interview Package	to set	Everything: all four volumes, the Practice Arena (Excel files, SQL, notebooks, datasets), and the assignments to work through on your own.
3	Projects	~₹150 each	Standalone, ready-made portfolio projects. Buy one, buy five.

The shape is right, and the third tier is the interesting one. A ₹500 book is a single transaction from each buyer. **₹150 projects are repeatable revenue from the same person**, they are small to build relative to a book, and they answer the question every Indian job-seeker actually has, which is not "what should I learn" but "**what do I put on my CV?**"

The thing that makes a ₹150 project worth buying

This is the one design decision that decides whether tier 3 works, so it is worth being blunt about.

GitHub is full of free code. If what we sell is a script that sends an email, we are competing with a thousand free repositories and we lose. Nobody needs another `smtplib` example.

What is genuinely scarce, and what each project must contain:

	Why it is the value
A realistic brief	Written the way a manager would ask, vague in the places managers are vague
Messy data	Not a clean CSV. Dates in two formats, a trailing space in a key, a duplicate, a missing region
The worked solution	Commented, runnable, with every output produced by running it — the book's standard
The reasoning	Why this approach and not the obvious one. This is Chapter 69A applied to a project
The interview narrative	The part nobody else sells. How to describe this project in ninety seconds, what the interviewer will probe, and the honest answer to "what would you do differently"

That last row is the product. A buyer is not paying ₹150 for code. They are paying for **something true to say in an interview**, and for the confidence of having actually run it.

Chapter 76A already teaches how to talk about a project and Chapter 82 has the take-home format, so the house style for that section already exists.

The catalogue, from Abhishek's list

Captured as specified, with what each one would contain and what to watch.

P1 · The daily report that sends itself

Raw data → processed → Excel report → emailed to a list.

The flagship, because it is the single most common real request a junior analyst gets, and because it is the Chapter 18 automation payoff as a standalone artifact.

Teaches	Reading messy raw data, aggregating, writing a formatted <code>.xlsx</code> , sending to a distribution list, scheduling it
Already in the book	Chapter 18 §18.15–18.18 (the report, the entry point, the <code>.bat</code> , Task Scheduler), Chapter 20 (delivery)
Interview narrative	"I automated a two-day manual month-end into a twenty-minute run" — the single most reusable line a junior analyst can have

The risk to design around: Outlook. Driving Outlook needs Windows *and* an installed, configured Outlook, which excludes Mac users, Linux users, anyone on Google Workspace, and most students. That is a large share of the buyers, and every one of them becomes a support email.

Recommendation: build it **SMTP-first** — which works everywhere, including Gmail with an app password — and ship the Outlook version as a clearly-labelled variant for people in a Microsoft office.

Same project, two delivery modules, one extra file. This is worth getting right because P1 is the one you will sell most.

P2 · Excel dashboards, ready-made and explained

A proper, good-looking dashboard, and how it was built.

Teaches	Model → pivots → charts → layout → interactivity, in the order you would build it
Already in the book	Chapter 11 (the spreadsheet mastered), Chapter 15 (what to plot), Chapter 70 §70.8 (dashboard critiques)
The hook	Most people have never seen a <i>well-built</i> dashboard file opened up. Shipping one that looks professional, with the build explained step by step, is genuinely rare

Ship the finished workbook **and** a stripped version so the buyer can build it themselves. The value is the gap between the two.

P3 · Python dashboards

The same, in code — Streamlit or Plotly Dash. A different audience from P2 and a different CV line.

P4 · The data-cleaning series

Jumbled data in, trustworthy data out. Abhishek is right that this is "an entirely different saleable project" — and it may be better as a *series* of three or four small ones than one large one, because each is a separate purchase.

Teaches	Profiling, duplicates, inconsistent categories, dates and time zones, validation rules that must return zero
Already in the book	Chapter 14 is the whole method, and Chapter 77 §77.8 now has the traps that bite on the way in
The hook	Every job ad says "data cleaning". Almost no portfolio shows it, because clean data is boring to publish. A project built on genuinely nasty data stands out

The Riverstone messy files already exist — `month_end_pack_2025_messy.xlsx`, `sales_export_with_problems.csv` — so the raw material is in the repository.

P5 · Logistics

Abhishek named this; the shape needs deciding. Likely candidates: delivery-route and cost analysis, inventory and stockout analysis, or supplier performance. Riverstone already has orders, products, suppliers and demand data, so it fits the house example without inventing a new company.

What else the catalogue should probably include

Three gaps that follow from what the book already contains, offered for a decision rather than assumed:

- **A SQL-only project.** A reporting layer built as a set of reviewed queries. Cheapest to build and support of anything on this list, because nothing has to run on the buyer's machine.
- **An end-to-end pipeline.** Ingest, clean, load, schedule, alert — the data-engineering CV line.
- **A first ML project done honestly.** With the baseline, the leakage check and the "would I ship this" verdict, which is what Chapter 74 argues for and what almost no portfolio project does.

Three things to settle before building any of them

1. The support boundary, written on the sales page. At ₹150 the net is about ₹131 after platform fees. **One support email per sale wipes out the margin.** The answer is not to refuse support, it is to design it out: pin every version, ship a sample dataset so nothing depends on the buyer's data, test on a clean machine, and state plainly what is and is not covered.

2. How a project differs from the Practice Arena. The Package (tier 2) already contains 643 practice files. If a ₹150 project looks like more of the same, a Package buyer will feel double-charged. The line that holds: **the Arena teaches the book; a project is a finished piece of work with a brief, messy data and an interview story.** Different purpose, different format. Say it on the page.

3. The projects Abhishek already has. You mentioned having a few of your own. Those are potentially the best items in the catalogue, because they are real work rather than teaching examples — and real beats realistic. **Send them over and I will tell you honestly which stand up as saleable projects and what each would need.**

Sequence

Agreed with the earlier instruction — the book first, then the projects.

Now	Finish the Interview Book: ~330 more questions to about 1,160, then the cover and back cover
Then	P1, built SMTP-first, as the template for the rest. Get one project completely right and the others follow a pattern
Then	P4, the cleaning series, since the messy data already exists
In parallel	Decide the Package price, and the support boundary